

House Price Prediction

X1= the transaction date (for example, 2013.250=2013 March, 2013.500=2013 June, etc.)

X2= the house age (unit: year)

X3= the distance to the nearest MRT station (unit: meter)

X4= the number of convenience stores in the living circle on foot (integer)

X5= the geographic coordinate, latitude. (unit: degree)

X6= the geographic coordinate, longitude. (unit: degree)

Y= house price of unit area (10000 New Taiwan Dollar/Ping, where Ping is a local unit, 1 Ping = 3.3 meter squared)

Data Pre-processing

Load data

In [1]:

```
# Import relevant packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
# Loading the data
housing = pd.read_excel('real_estate_valuation.xlsx')

# Converting to dataframe
df_housing = pd.DataFrame(housing)

# Making a copy of the original dataframe to work on
df = housing.copy(deep=True)
```

In [3]:

```
# previewing the first 5 rows
df.head()
```

Out[3]:

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
0	1	2012.916667	32.0	84.87882	10	24.98298	121.54024	37.9
1	2	2012.916667	19.5	306.59470	9	24.98034	121.53951	42.2
2	3	2013.583333	13.3	561.98450	5	24.98746	121.54391	47.3
3	4	2013.500000	13.3	561.98450	5	24.98746	121.54391	54.8
4	5	2012.833333	5.0	390.56840	5	24.97937	121.54245	43.1

In [4]:

```
# Previewing the last 5 rows
df.tail()
```

Out[4]:

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
409	410	2013.000000	13.7	4082.01500	0	24.94155	121.50381	15.4
410	411	2012.666667	5.6	90.45606	9	24.97433	121.54310	50.0
411	412	2013.250000	18.8	390.96960	7	24.97923	121.53986	40.6
412	413	2013.000000	8.1	104.81010	5	24.96674	121.54067	52.5
413	414	2013.500000	6.5	90.45606	9	24.97433	121.54310	63.9

In [5]:

```
# preview the number of attributes
df.shape[1]
```

Out[5]:

8

In [6]:

```
# Preview the column names  
df.columns
```

Out[6]:

```
Index(['No', 'X1 transaction date', 'X2 house age',  
      'X3 distance to the nearest MRT station',  
      'X4 number of convenience stores', 'X5 latitude', 'X6 longitud  
e',  
      'Y house price of unit area'],  
      dtype='object')
```

In [7]:

```
# Checking for missing values  
df.isna().sum()
```

Out[7]:

```
No                                0  
X1 transaction date              0  
X2 house age                    0  
X3 distance to the nearest MRT station 0  
X4 number of convenience stores  0  
X5 latitude                     0  
X6 longitude                    0  
Y house price of unit area       0  
dtype: int64
```

In [8]:

```
# No missing values within this dataset, no imputer required.
```

In [9]:

```
# checking the datatypes of each column  
df.dtypes
```

Out[9]:

```
No                                int64  
X1 transaction date              float64  
X2 house age                    float64  
X3 distance to the nearest MRT station float64  
X4 number of convenience stores  int64  
X5 latitude                     float64  
X6 longitude                    float64  
Y house price of unit area       float64  
dtype: object
```

In [10]:

```
# Summarising the dataframe
df.describe()
```

Out[10]:

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude
count	414.000000	414.000000	414.000000	414.000000	414.000000	414.000000	414.000000
mean	207.500000	2013.148953	17.712560	1083.885689	4.094203	24.969030	121.533361
std	119.655756	0.281995	11.392485	1262.109595	2.945562	0.012410	0.015347
min	1.000000	2012.666667	0.000000	23.382840	0.000000	24.932070	121.473530
25%	104.250000	2012.916667	9.025000	289.324800	1.000000	24.963000	121.528085
50%	207.500000	2013.166667	16.100000	492.231300	4.000000	24.971100	121.538630
75%	310.750000	2013.416667	28.150000	1454.279000	6.000000	24.977455	121.543305
max	414.000000	2013.583333	43.800000	6488.021000	10.000000	25.014590	121.566270

Cleaning the data

In [11]:

```
# removing irrelevant column
df.drop(columns = 'No', inplace = True)
```

In [12]:

```
# Cleaning the column names
df.rename(columns = {'X1 transaction date': 'transaction_date',
                    'X2 house age': 'house_age',
                    'X3 distance to the nearest MRT station': 'mrt_distance',
                    'X4 number of convenience stores': 'no_of_stores',
                    'X5 latitude': 'latitude',
                    'X6 longitude': 'longitude',
                    'Y house price of unit area': 'house_price'}, inplace = True)
```

In [13]:

```
df.head()
```

Out[13]:

	transaction_date	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
0	2012.916667	32.0	84.87882	10	24.98298	121.54024	37.9
1	2012.916667	19.5	306.59470	9	24.98034	121.53951	42.2
2	2013.583333	13.3	561.98450	5	24.98746	121.54391	47.3
3	2013.500000	13.3	561.98450	5	24.98746	121.54391	54.8
4	2012.833333	5.0	390.56840	5	24.97937	121.54245	43.1

Preliminary Analysis

In [14]:

```
# What's the lowest house price  
print(f"The lowest house price is {round(df.house_price.min(),2)}/m\u00b2")
```

The lowest house price is 7.6/m²

In [15]:

```
print(f"The highest house price is {round(df.house_price.max(),2)}/m\u00b2")
```

The highest house price is 117.5/m²

In [16]:

```
print(f"The average house price is {round(df.house_price.mean(),2)}/m\u00b2")
```

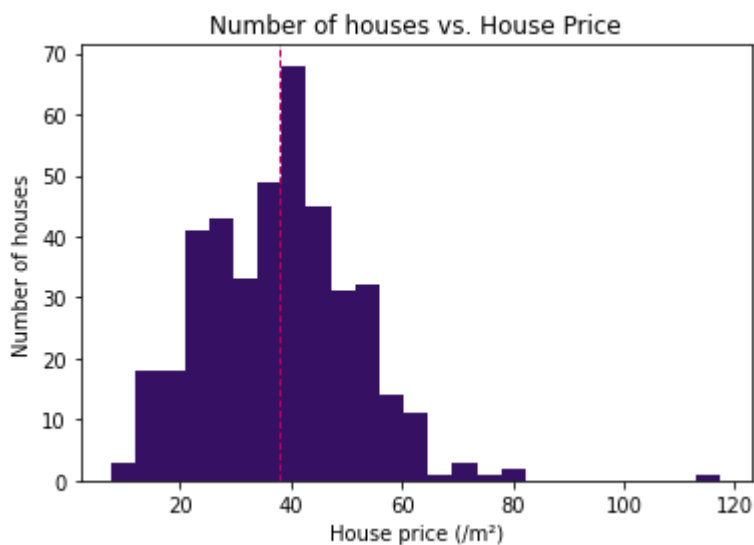
The average house price is 37.98/m²

In [17]:

```
# Checking the distribution of the target variable (house_price)
plt.hist(df["house_price"], bins=25, color="#361163")
plt.xlabel("House price (/m\u00b2)")
plt.ylabel("Number of houses")
plt.title("Number of houses vs. House Price")
plt.ticklabel_format(style='plain')

# Plotting the line for average house price (house_price)
plt.axvline(df["house_price"].mean(), color='#B70062', linestyle='dashed', linewidth=2)
plt.show()

# ===== ADD COMMA TO VALUES
# ===== EXPAND THE TICKS?
```



This plot shows the house price (house_price) and the number of houses sold in the price range. It is skewed right, where the right side is the highest house price sold, with only one houses selling for 117.5/ m^2 . Most houses sold between 10/ m^2 - 65/ m^2 . Around 15 houses' sale price was between 60/ m^2 - 80/ m^2 .

In [18]:

```
# What's the average distance to the nearest MRT station in meters?
print("The average distance to the nearest MRT station is",
      round(df.mrt_distance.mean(),2),
      "meters")
```

The average distance to the nearest MRT station is 1083.89 meters

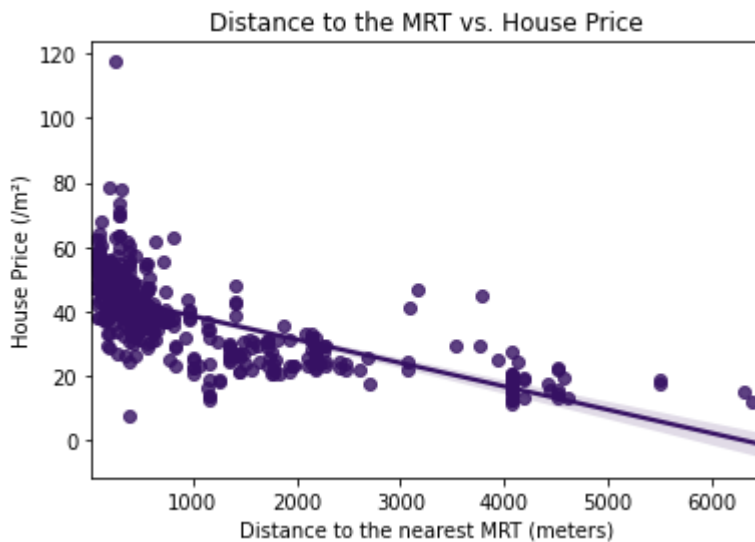
In [19]:

```
# Does the distance to the nearest MRT have an impact on the house price?
sns.regplot(x="mrt_distance", y="house_price", data=df, color="#361163")
plt.ticklabel_format(style='plain')
plt.title('Distance to the MRT vs. House Price')
plt.xlabel('Distance to the nearest MRT (meters)')
plt.ylabel('House Price (/m²)')

# ===== HIGHLIGHT THE OUTLIERS
```

Out[19]:

Text(0, 0.5, 'House Price (/m²)')



In [20]:

```
# Investigating outlier 1
df[df['house_price'] == df.house_price.max()]
```

Out[20]:

	transaction_date	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
270	2013.333333	10.8	252.5822	1	24.9746	121.53046	117.5

In [21]:

```
# Investigating outlier 2
df[df['house_price'] == df.house_price.min()]
```

Out[21]:

	transaction_date	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
113	2013.333333	14.8	393.2606	6	24.96172	121.53812	7.6

- It seems that the houses with the shortest distance to the MRT station have a higher house price.
- It is also evident that the houses with a shorter distance to the MRT station, sold more than those that are further away from the nearest MRT station.
- The outlier on the plot shows the house with highest price overall sold at 117.5/m², situated 252.58 meters away to the nearest MRT station.
- However, another outlier shows that the house sold at the lowest price was situated close to an MRT station with only 393.26 meters away.
- Nevertheless, the distance to the nearest MRT station (mrt_distance) is likely to be a strong predictor for the house price (house_price)

In [22]:

```
# Removing outliers to make to prevent skewed results
df.drop(labels=270, axis=0, inplace = True)
df.drop(labels=113, axis=0, inplace = True)
```

In [23]:

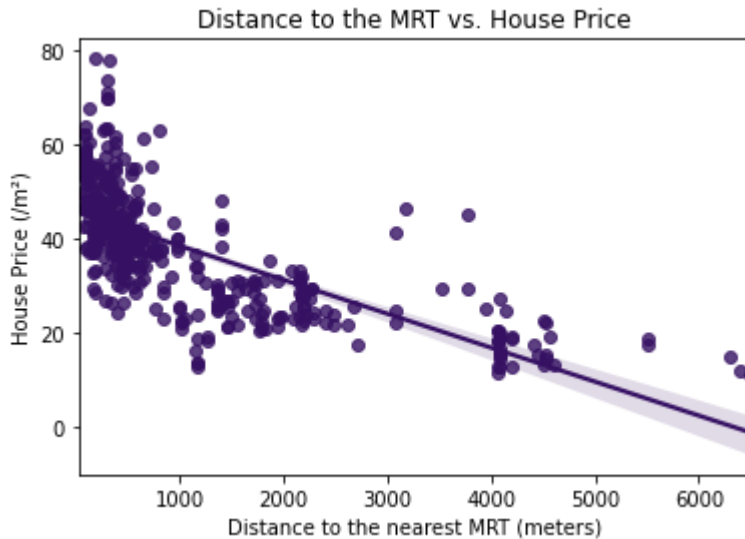
```
# Checking the outliers have been dropped
df.shape[0]
```

Out[23]:

412

In [24]:

```
# Rechecking the previous plot
sns.regplot(x="mrt_distance", y="house_price", data=df, color="#361163")
plt.title('Distance to the MRT vs. House Price')
plt.xlabel('Distance to the nearest MRT (meters)')
plt.ylabel('House Price (/m²)')
plt.show()
```

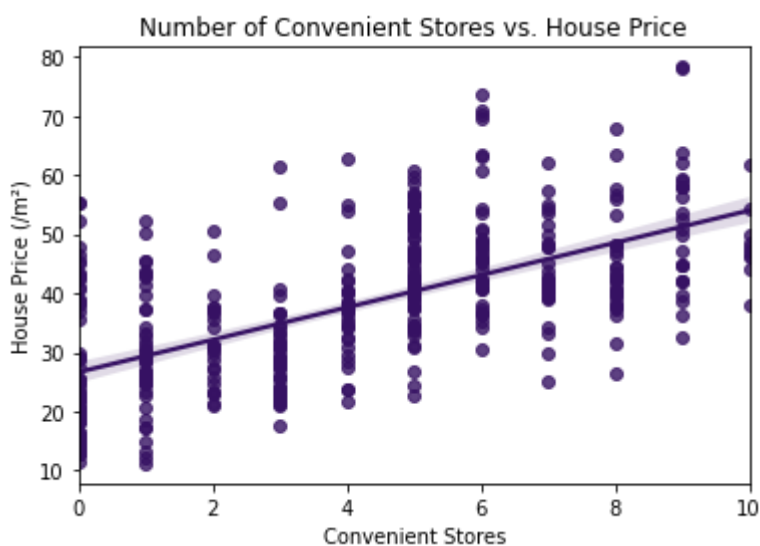


In [25]:

```
# Does having more convenient stores near the living circle increase the house price
sns.regplot(x="no_of_stores", y="house_price", data=df, color="#361163")
plt.title('Number of Convenient Stores vs. House Price')
plt.xlabel('Convenient Stores')
plt.ylabel('House Price (/m²)')
```

Out[25]:

```
Text(0, 0.5, 'House Price (/m²)')
```



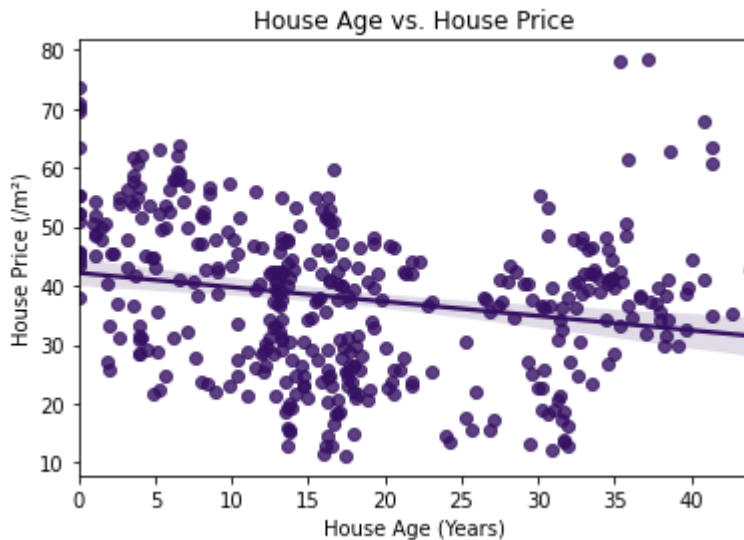
- This plot shows the number of the convenient stores in the living circle (no_of_stores) has a positive linear trend and could be a good predictor for the house price (house_price)

In [26]:

```
# Does the age of the house have impact on the sale price?
sns.regplot(x="house_age", y="house_price", data=df, color="#361163")
plt.title('House Age vs. House Price')
plt.xlabel('House Age (Years)')
plt.ylabel('House Price (/m\u00b2)')
```

Out[26]:

Text(0, 0.5, 'House Price (/m²)')

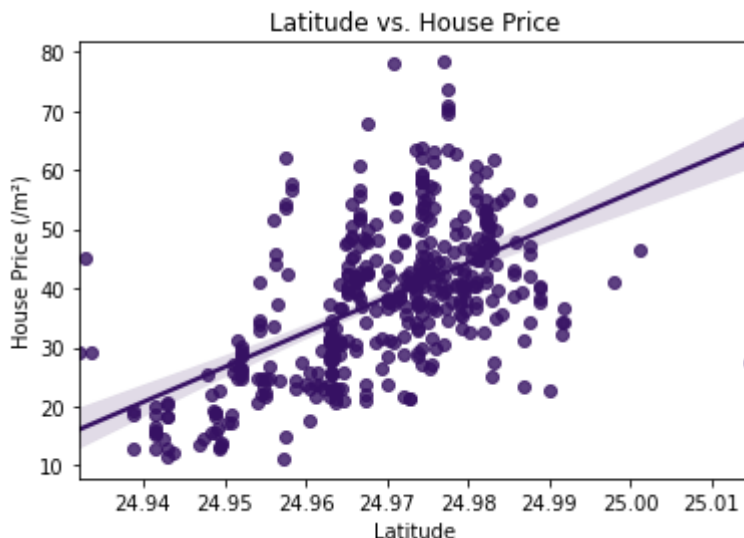


In [27]:

```
# Does latitude have impact on the sale price?
sns.regplot(x="latitude", y="house_price", data=df, color="#361163")
plt.title('Latitude vs. House Price')
plt.xlabel('Latitude')
plt.ylabel('House Price (/m\u00b2)')
```

Out[27]:

Text(0, 0.5, 'House Price (/m²)')

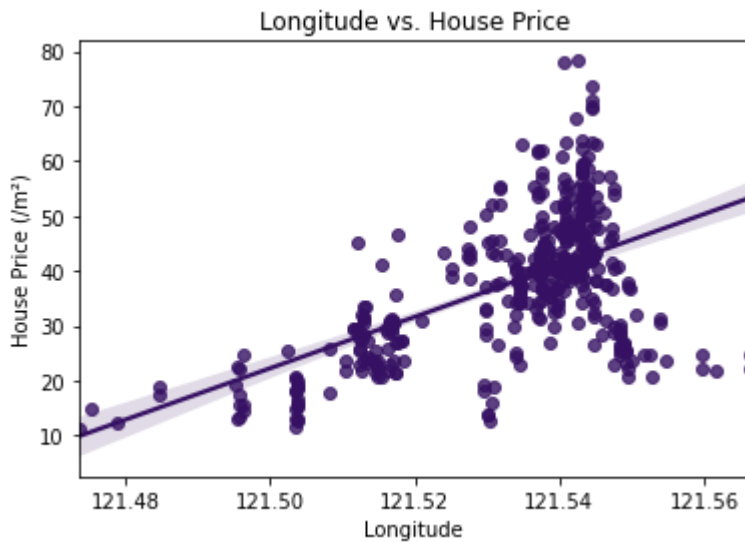


In [28]:

```
# Does longitude have impact on the sale price?
sns.regplot(x="longitude", y="house_price", data=df, color="#361163")
plt.title('Longitude vs. House Price')
plt.xlabel('Longitude')
plt.ylabel('House Price (/m\u00b2)')
```

Out[28]:

Text(0, 0.5, 'House Price (/m²)')

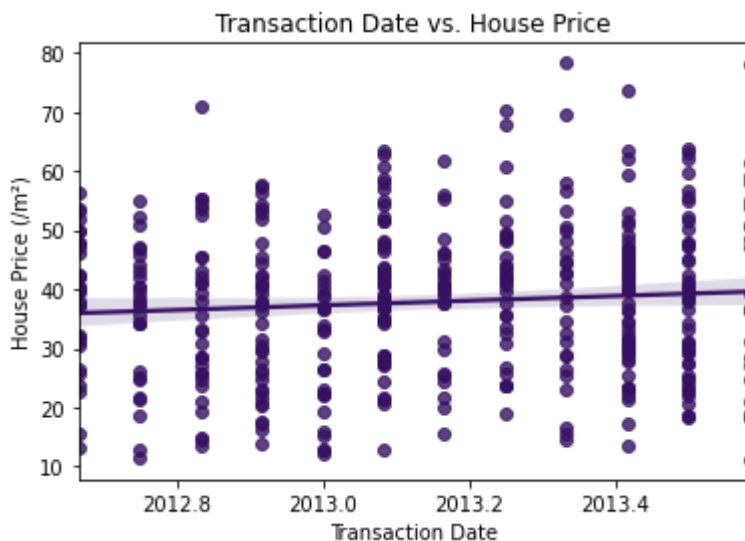


In [29]:

```
# How does the transaction date have impact on the sale price?
sns.regplot(x="transaction_date", y="house_price", data=df, color="#361163")
plt.title('Transaction Date vs. House Price')
plt.xlabel('Transaction Date')
plt.ylabel('House Price (/m\u00b2)')
```

Out[29]:

Text(0, 0.5, 'House Price (/m²)')



In [30]:

df.head()

Out[30]:

	transaction_date	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
0	2012.916667	32.0	84.87882	10	24.98298	121.54024	37.9
1	2012.916667	19.5	306.59470	9	24.98034	121.53951	42.2
2	2013.583333	13.3	561.98450	5	24.98746	121.54391	47.3
3	2013.500000	13.3	561.98450	5	24.98746	121.54391	54.8
4	2012.833333	5.0	390.56840	5	24.97937	121.54245	43.1

In [31]:

```
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

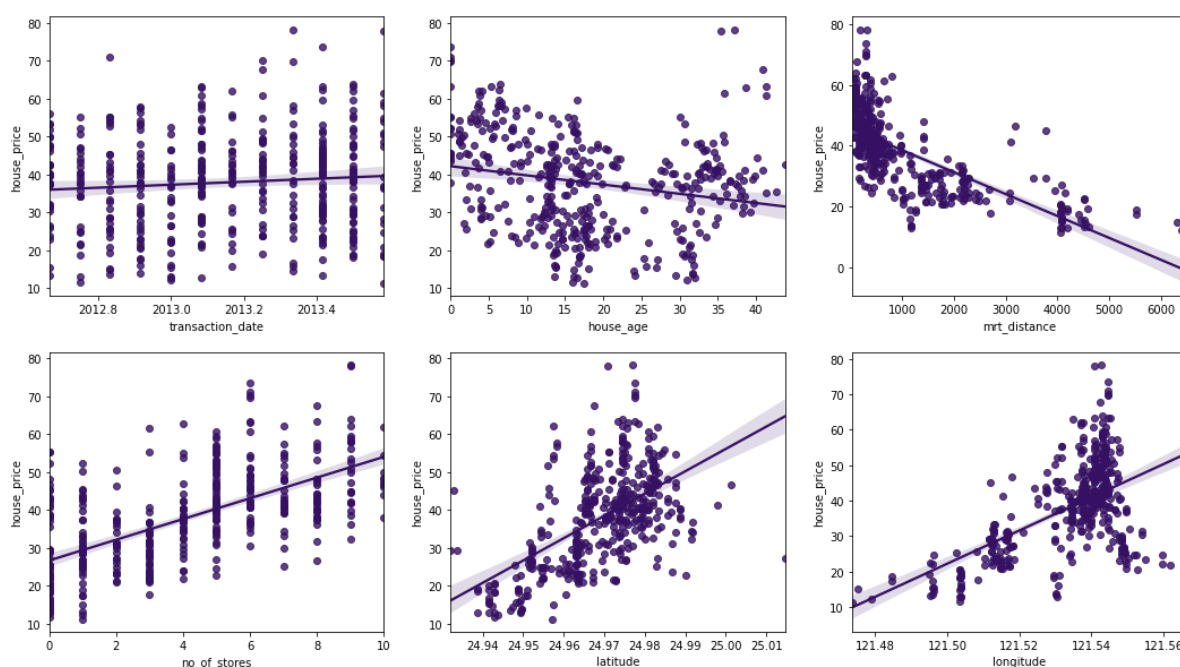
fig.suptitle('Relationships between the different attributes and the target variable')

sns.regplot(ax=axes[0, 0], data=df, x="transaction_date", y="house_price", color="#361163")
sns.regplot(ax=axes[0, 1], data=df, x='house_age', y='house_price', color="#361163")
sns.regplot(ax=axes[0, 2], data=df, x='mrt_distance', y='house_price', color="#361163")
sns.regplot(ax=axes[1, 0], data=df, x='no_of_stores', y='house_price', color="#361163")
sns.regplot(ax=axes[1, 1], data=df, x='latitude', y='house_price', color="#361163")
sns.regplot(ax=axes[1, 2], data=df, x='longitude', y='house_price', color="#361163")
```

Out[31]:

<AxesSubplot:xlabel='longitude', ylabel='house_price'>

Relationships between the different attributes and the target variable (house_price)

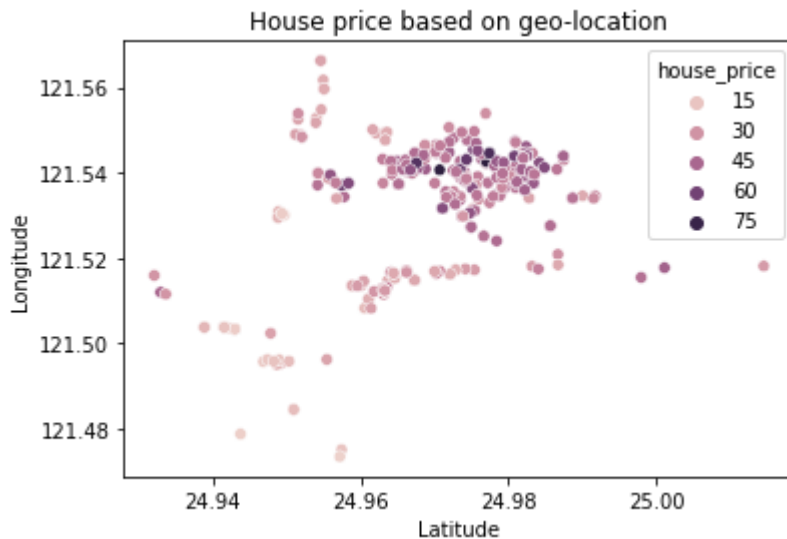


In [32]:

```
# Latitude, Longitude and House Price
sns.scatterplot(data=df, x="latitude", y="longitude", hue="house_price", sizes=(20,200))
plt.xlabel("Latitude")
plt.ylabel("Longitude")
plt.title("House price based on geo-location")
```

Out[32]:

Text(0.5, 1.0, 'House price based on geo-location')



Feature Selection

We can see that the date of the house sale (transaction_date) has little to no correlation with the sale price (house_price). Therefore, I am removing this column.

In [33]:

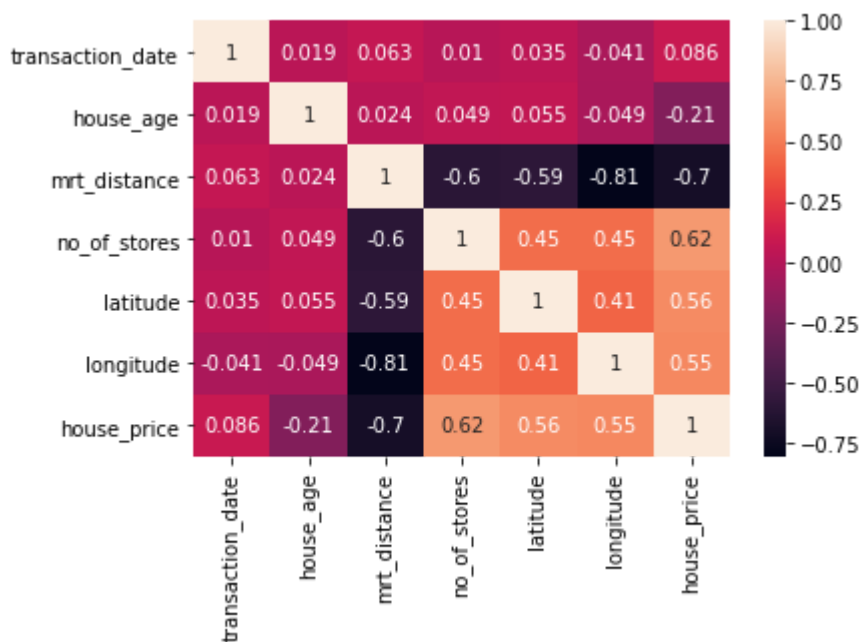
```
# Filter features by correlation
# What are the correlation of each attribute to the house price (house_price)?
print(abs(df.corr()[ 'house_price' ]))

# Visualisation
sns.heatmap(df.corr(),annot=True)
```

```
transaction_date    0.086052
house_age           0.213846
mrt_distance        0.701939
no_of_stores        0.620631
latitude            0.564572
longitude           0.554680
house_price         1.000000
Name: house_price, dtype: float64
```

Out[33]:

<AxesSubplot:>



In [34]:

```
col_drop = df[ ['transaction_date' ] ]
df.drop(columns = col_drop, inplace = True)
```

In [35]:

```
df.head()
```

Out[35]:

	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
0	32.0	84.87882	10	24.98298	121.54024	37.9
1	19.5	306.59470	9	24.98034	121.53951	42.2
2	13.3	561.98450	5	24.98746	121.54391	47.3
3	13.3	561.98450	5	24.98746	121.54391	54.8
4	5.0	390.56840	5	24.97937	121.54245	43.1

Feature Scaling

Code adapted from: <https://towardsdatascience.com/feature-scaling-effectively-choose-input-variables-based-on-distributions-3032207c921f> (<https://towardsdatascience.com/feature-scaling-effectively-choose-input-variables-based-on-distributions-3032207c921f>) <https://www.baeldung.com/cs/normalization-vs-standardization> (<https://www.baeldung.com/cs/normalization-vs-standardization>)

In [36]:

```
# Import relevant package for feature scaling
from sklearn.preprocessing import MinMaxScaler
```

In [37]:

```
# Normalized dataset
min_max = MinMaxScaler()
df_norm = min_max.fit_transform(df)
df_norm = pd.DataFrame(df_norm, columns = df.columns)
df_norm.head()
```

Out[37]:

	house_age	mrt_distance	no_of_stores	latitude	longitude	house_price
0	0.730594	0.009513	1.0	0.616941	0.719323	0.397914
1	0.445205	0.043809	0.9	0.584949	0.711451	0.461997
2	0.303653	0.083315	0.5	0.671231	0.758896	0.538003
3	0.303653	0.083315	0.5	0.671231	0.758896	0.649776
4	0.114155	0.056799	0.5	0.573194	0.743153	0.475410

Model Selection

In [38]:

```
# Import relevant packages for building the models
np.random.seed(42)
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor

# Import relevant packages for the performance evaluation
from sklearn.metrics import r2_score, explained_variance_score, mean_absolute_error
from sklearn.metrics import mean_squared_error, median_absolute_error
# from sklearn.model_selection import cross_val_score ===== DELETE THIS
```

Linear Regression

Fitting the model and making predictions

In [39]:

```
# Splitting features(X) and target(y)
X = df_norm.drop("house_price", axis = 1)
y = df_norm.house_price

# Splitting into training and testing sets
np.random.seed(42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# Build the Linear Regression Model on original dataset
lr = LinearRegression()

# Fit the Linear Regression Model
lr.fit(X_train, y_train)

# Using the model to make a prediction
y_pred_lr = lr.predict(X_test)
```

In [40]:

```
# Checking the shapes of training and testing data split
print(X_train.shape, y_train.shape,
      X_test.shape, y_test.shape)
```

```
(329, 5) (329,) (83, 5) (83,)
```

Performance Evaluation

In [41]:

```

print("The performance evaluation of the Linear Regression model with the original dataset are as follows:")

# Checking the R squared score of the model [Best score: 1]
print("R-squared score is:", round(r2_score(y_test, y_pred_lr),2))

# Checking the Mean Squared Error of the model (MSE) [Best Value: 0.0]
print("Mean squared error is:", round(mean_squared_error(y_test, y_pred_lr),2))

# Checking the Mean Absolute Error (MAE), which shows the average difference between predicted and actual values
print("Mean absolute error is:", round(mean_absolute_error(y_test, y_pred_lr),2))

# ===== DELETE BELOW
# # Checking the Median Absolute Error (MAD) [Best value: 0.0]
# print("Median absolute error is:", median_absolute_error(y_test, y_pred_lr))

# # Checking the Explained Variance Score
# # (Best possible score: 1)
# print("Explained variance score is:", explained_variance_score(y_test, y_pred_lr))

```

The performance evaluation of the Linear Regression model with the original dataset are as follows:

R-squared score is: 0.67

Mean squared error is: 0.01

Mean absolute error is: 0.08

Visualizing the model

In [42]:

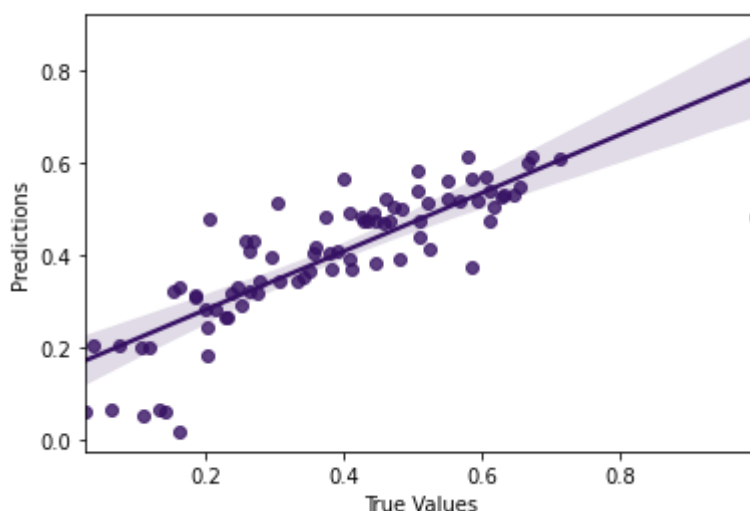
```

sns.regplot(x=y_test, y=y_pred_lr, color="#361163")
plt.xlabel("True Values")
plt.ylabel("Predictions")

```

Out[42]:

Text(0, 0.5, 'Predictions')



Random Forest Regressor

Fitting a 'Random Forest Regressor' model

In [43]:

```
from sklearn.ensemble import RandomForestRegressor
```

In [44]:

```
# Splitting features(X) and target(y)
X = df_norm.drop("house_price", axis = 1)
y = df_norm.house_price

# Splitting into training and testing sets
np.random.seed(42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# Build the SVR Model
rf = RandomForestRegressor()

# Fit the model
rf.fit(X_train, y_train)

# Using the model to make a prediction
y_pred_rf = rf.predict(X_test)
```

Performance Evaluation

In [45]:

```
print("The performance evaluation of the Random Forest Regression model are as follows")

# Checking the R squared score of the model
# (Best score: 1)
print("R-squared score is:", round(r2_score(y_test, y_pred_rf),2))

# Checking the Mean Squared Error of the model (MSE)
# (Best Value: 0.0)
print("Mean squared error is:", round(mean_squared_error(y_test, y_pred_rf),2))

# Checking the Mean Absolute Error (MAE), which shows the average difference between
# Best value: 0.0
print("Mean absolute error is:", round(mean_absolute_error(y_test, y_pred_rf),2))
```

The performance evaluation of the Random Forest Regression model are as follows:

```
R-squared score is: 0.71
Mean squared error is: 0.01
Mean absolute error is: 0.07
```

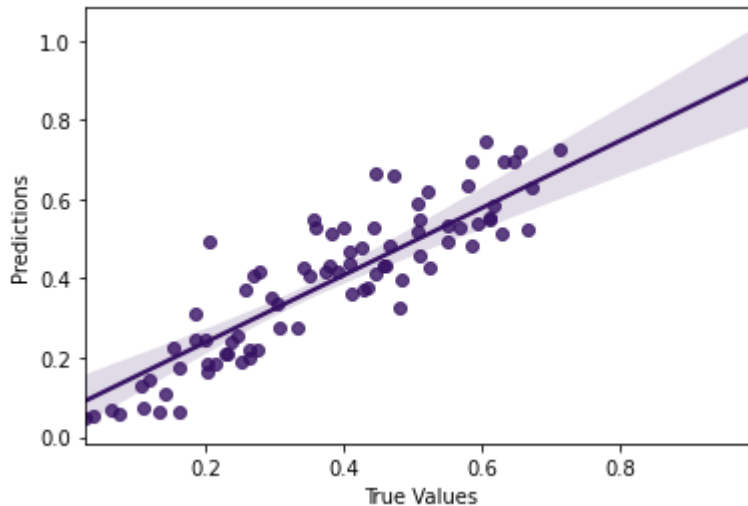
Visualize the model

In [46]:

```
sns.regplot(x=y_test, y=y_pred_rf, color="#361163")
plt.xlabel("True Values")
plt.ylabel("Predictions")
```

Out[46]:

Text(0, 0.5, 'Predictions')



KNearestNeighbor

Fitting a 'K-nearest Neighbor' model

In [48]:

```
# Splitting features(X) and target(y)
X = df_norm.drop("house_price", axis = 1)
y = df_norm.house_price

# Splitting into training and testing sets
np.random.seed(42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

# Build the K-nearest Regressor model on original dataset
knr = KNeighborsRegressor(n_neighbors=5)

# Fit the K-nearest Regressor model
knr.fit(X_train, y_train)

# Using the model to make a prediction
y_pred_knr = knr.predict(X_test)
```

Performance Evaluation

In [49]:

```
print("The performance evaluation of the K-nearest Regressor model are as follows:\n")

# Checking the R squared score of the model
# (Best score: 1)
print("R-squared score is:", round(r2_score(y_test, y_pred_knr),2))

# Checking the Mean Squared Error of the model (MSE)
# (Best Value: 0.0)
print("Mean squared error is:", round(mean_squared_error(y_test, y_pred_knr),2))

# Checking the Mean Absolute Error (MAE), which shows the average difference between
# Best value: 0.0
print("Mean absolute error is:", round(mean_absolute_error(y_test, y_pred_knr),2))
```

The performance evaluation of the K-nearest Regressor model are as follows:

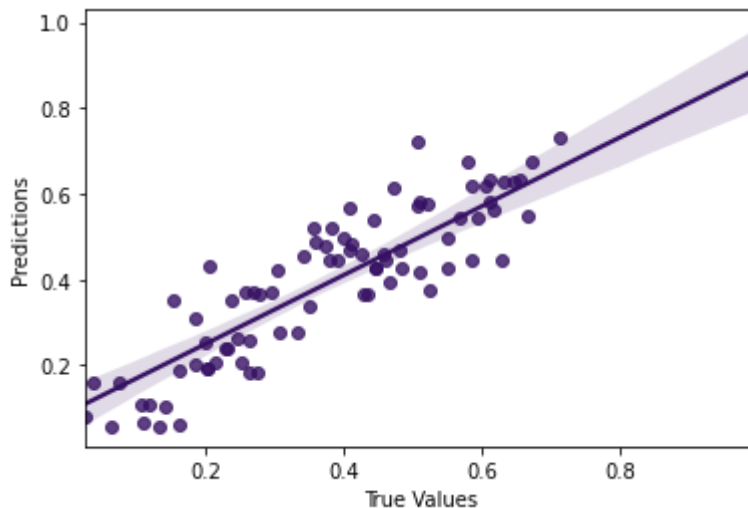
R-squared score is: 0.73
 Mean squared error is: 0.01
 Mean absolute error is: 0.07

In [50]:

```
sns.regplot(x=y_test, y=y_pred_knr, color="#361163")
plt.xlabel("True Values")
plt.ylabel("Predictions")
```

Out[50]:

Text(0, 0.5, 'Predictions')



Improving the model

K-Fold Cross Validation

K-folds

Leave One Out

Computationally expensive but suitable for small datasets

<https://medium.com/analytics-vidhya/using-cross-validation-to-evaluate-different-models-regression-5f61ec89531> (<https://medium.com/analytics-vidhya/using-cross-validation-to-evaluate-different-models-regression-5f61ec89531>)

In [52]:

```
# Importing relevant package
# from sklearn.model_selection import LeaveOneOut
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn import metrics

# Linear Regression
lr_cv = round(np.mean(cross_val_score(lr, X, y, cv=5))*100,2)
print(f"Accuracy of Linear Regression with Cross Validation: {lr_cv}%")

# Random Forest Regressor
rf_cv = round(np.mean(cross_val_score(rf, X, y, cv=5))*100,2)
print(f"Accuracy of Random Forest Regressor with Cross Validation: {rf_cv}%")

# K-Nearest Regressor
knr_cv = round(np.mean(cross_val_score(knr, X, y, cv=5))*100,2)
print(f"Accuracy of K-Nearest Regressor with Cross Validation: {knr_cv} %")
```

Accuracy of Linear Regression with Cross Validation: 61.47%
Accuracy of Random Forest Regressor with Cross Validation: 73.69%
Accuracy of K-Nearest Regressor with Cross Validation: 71.27 %

In []: